

Name: Sahal Shrestha

Date: 30/01/2026

Section A:

- 1) B
- 2) C
- 3) B
- 4) B
- 5) B
- 6) A
- 7) C
- 8) C
- 9) D
- 10) B
- 11) B
- 12) B
- 13) B
- 14) B
- 15) A
- 16) B
- 17) B
- 18) B
- 19) B
- 20) B

Section B

Link: [Click here](#)

1.

1. Chunking Strategy

To ensure high retrieval accuracy, we used a **Recursive Content-Aware** chunking strategy.

Implementation Details

- **Tool:** LangChain's `RecursiveCharacterTextSplitter`.
- **Chunk Size:** `500` characters. This provides a balance between sufficient context for the model and focused retrieval.
- **Chunk Overlap:** `50` characters. This ensures that information spanning across chunks maintains continuity.
- **Separators:** `["\n\n", "\n", " ", ""]`
 - **Priority 1** (`\n\n`): We first attempt to split by double newlines to keep individual product entries as atomic units.
 - **Priority 2** (`\n`): If a product entry is too long, we split by line breaks.
 - **Priority 3** (): As a final resort, we split at word boundaries.

2. Embedding Configuration

Embeddings convert text into high-dimensional vectors that the chatbot uses to find relevant information.

Implementation Details

- **Model:** `models/embedding-001` via **Google Generative AI (Gemini)**.
- **Provider:** `langchain-google-genai`.
- **Vector Store:** **ChromaDB**.
- **Persistence:** Data is persisted locally in the `/home/labuser/assessment/chroma_db` directory.

3. Workflow Integration

During a query:

1. The user's question is embedded using the same `embedding-001` model.
2. A **Similarity Search** is performed against the ChromaDB vector store.
3. The top `k=3` most relevant chunks are retrieved and injected into the LLM prompt as context.

2.

Walkthrough - Product Knowledge Base RAG

I have completed the end-to-end implementation of the Product Knowledge Base RAG system.

Changes Made

1. Knowledge Base Ingestion
 - Developed `ingest_data.py` to process product details.
 - Used `RecursiveCharacterTextSplitter` for chunking.
 - Stored embeddings in a persistent `ChromaDB` local instance.
2. RAG Inference Chain
 - Created `rag_chain.py` using LangChain Expression Language (LCEL).
 - **Model:** Configured `gemini-2.5-flash` as requested.
 - **Retriever:** Connected to the ChromaDB vector store.
 - **Prompting:** Implemented a system prompt that enforces answering strictly based on the provided context.

3.

1. LangGraph Workflow Structure

Developed `langgraph_workflow.py` which implements the following nodes:

- **Node 1: Classifier:** Uses `gemini-2.5-flash` to categorize queries into `products`, `returns`, or `general`.
- **Node 2: RAG Responder:** Reuses our RAG logic to answer product-specific questions when the classifier detects a product query.
- **Node 3: Escalation:** Intercepts `returns` or `general` queries and provides a structured escalation response.

2. Conditional Routing

Implemented routing logic that directs the flow dynamically:

- `products` → **RAG Responder**
- `returns` or `general` → **Escalation**

4.

The screenshot shows a web browser window displaying the API documentation for the Product Knowledge Base Chatbot. The URL in the address bar is `localhost:8000/docs#/`. The title of the page is "Product Knowledge Base Chatbot" with version `0.1.0` and `OAS 3.1`. Below the title is a link to `/openapi.json`.

The main content area is titled "default". It contains two sections:

- GET / Root**: A blue button representing a GET request to the root endpoint.
- POST /query Query Chatbot**: A green button representing a POST request to the `/query` endpoint, which is described as "Query Chatbot".

Below these sections is a "Schemas" section containing a link to `HTTPValidationError > Expand all object`.

At the bottom, there is a table titled "Code Details" showing the response for a 200 status code:

Code	Details
200	<p>Response body</p> <pre>{ "question": "what is the price of Smartwatch Pro X", "category": "products", "answer": "The price of SmartWatch Pro X is ₹15,999." }</pre>

System Architecture

The system is built with a layered architecture for reliability and intelligence:

- Ingestion Layer**: Processes `product_details.txt` into high-quality vector embeddings stored in a local `ChromaDB`.
- Orchestration Layer**: A `LangGraph workflow` that classifies user intent and routes queries to either the RAG Responder or an Escalation handler.
- Service Layer**: A `FastAPI application` that provides a robust web interface for user interaction.

Outputs:



